

Optimization in Python

Kevin Carlberg (Meta)

August 5, 2022

Optimization tools in Python

We will go over and use two tools:

1. `scipy.optimize`
2. CVXPY

See `quadratic_minimization.ipynb`

- ▶ User inputs defined in the second cell
- ▶ Enables exploration of how problem attributes affect optimization-solver performance

`scipy.optimize`

Outline

`scipy.optimize`

CVXPY

Example: `quadratic_minimization.ipynb`

`scipy.optimize`

scipy.optimize


scipy.optimize: sub-package of SciPy, which is an open source Python library for scientific computing

- ▶ Analogous to Matlab's optimization toolbox
- ▶ Capabilities
 - ▶ Optimization
 - ▶ **Local optimization**
 - ▶ Equation minimizers
 - ▶ **Global optimization**
 - ▶ Fitting (nonlinear least squares)
 - ▶ Root finding
 - ▶ Linear Programming
 - ▶ Utilities (e.g., `check_grad` for verifying analytic gradients)

scipy.optimize interface

Requires the user to define a function in Python

- ▶ Can be **black box**: no closed-form mathematical expression needed!
- ▶ Only the function value $f(x)$ is required
- ▶ Can optionally provides the gradient $\nabla f(x)$ and Hessian $\nabla^2 f(x)$
- ▶ *Example*: evaluating f constitutes a run of a complicated simulation code
- ▶ **Drawback**: cannot exploit special structure underlying f



scipy.optimize: local optimization algorithms

Unconstrained minimization

- ▶ *Derivative free*: no gradient or Hessian
 - ▶ Nelder-Mead: simplex
 - ▶ Powell: sequential minimization along each vector in a direction set
- ▶ *Gradient-based*: gradient only (no Hessian)
 - ▶ CG: nonlinear conjugate gradient
 - ▶ BFGS: quasi-Newton BFGS method
- ▶ *Gradient-based*: gradient and Hessian can be specified
 - ▶ Newton-CG: approximately solves Newton system using CG (truncated Newton method)
 - ▶ dogleg: dog-leg trust-region algorithm. Hessian must be SPD
 - ▶ trust-ncg: Newton conjugate gradient trust-region method

scipy.optimize: local optimization algorithms

Constrained minimization (all are gradient-based)

- ▶ Only bound constraints
 - ▶ L-BFGS-B: limited memory BFGS bound constrained optimization
 - ▶ TNC: truncated Newton allows for upper and lower bounds
- ▶ General constraints
 - ▶ COBYLA: Constrained Optimization BY Linear Approximation
 - ▶ SLSQP: Sequential Least Squares Programming

scipy.optimize: global optimization algorithms

Global optimization (generally derivative free)

- ▶ `basinhopping`: stochastic algorithm by Wales and Doye,
 - ▶ useful when the function has many minima separated by large barriers
- ▶ `brute`: brute force minimization over a specified range
- ▶ `differential_evolution`: an evolutionary algorithm by Storn and Price
- ▶ `shgo`: simplicial homology global optimization (can specify gradient/Hessian)
- ▶ `dual_annealing`: generalized simulated annealing algorithm

CVXPY

Outline

`scipy.optimize`

CVXPY

Example: `quadratic_minimization.ipynb`

Modeling languages for convex optimization

- ▶ High-level language support for convex optimization has been developed recently
 1. Describe problem in high-level language
 2. Description automatically transformed to standard form
 3. Solved by standard solver, transformed back
- ▶ Implementations:
 - ▶ YALMIP, CVX (Matlab)
 - ▶ CVXPY (Python)
 - ▶ Convex.jl (Julia)
 - ▶ CVXR (R)
- ▶ Benefits:
 - ▶ Easy to perform rapid prototyping
 - ▶ Can exploit special structure because have *full mathematical description*
 - ▶ Let users focus on *what* their model should be instead of *how* to solve it
 - ▶ No algorithm tuning or babysitting
- ▶ Drawbacks:
 - ▶ Won't work if your problem isn't convex
 - ▶ Need explicit mathematical formulas for the objective and constraints (no black box!)

- ▶ **CVXPY**: “a Python-embedded **modeling language** for convex optimization problems. It allows you to **express your problem in a natural way** that follows the math, rather than in the **restrictive standard form** required by solvers.”

```
from cvxpy import *  
x = Variable(n)  
cost = sum_squares(A*x-b) + gamma*norm(x,1) # explicit formula!  
prob = Problem(Minimize(cost, [norm(x,"inf") <=1]))  
opt_val = prob.solve()  
solution = x.value
```

- ▶ solve method converts problem to standard form, solves and assigns opt_val attributes

CVXPY usage

- ▶ `cvxpy.Problem`: optimization problem
- ▶ `cvxpy.Variable`: optimiation variable
- ▶ `cvxpy.Minimize`: minimization function
- ▶ `cvxpy.Parameter`: symbolic representations of constants
 - ▶ can change the value of a constant without reconstructing the entire problem
 - ▶ can enforce to be positive or negative on construction
- ▶ Constraints simply Python lists
- ▶ Many functions implemented: see cvxpy.org website for list

Complete CVXPY example

```
import cvxpy as cvx
# Create two scalar optimization variables (CVXPY Variable)
x = cvx.Variable()
y = cvx.Variable()
# Create two constraints (Python list)
constraints = [x + y == 1, x - y >= 1]
# Form objective
obj = cvx.Minimize(cvx.square(x - y))
# Form and solve problem
prob = cvx.Problem(obj, constraints)
prob.solve() # Returns the optimal value.
print("status:", prob.status)
print("optimal value", prob.value)
print("optimal var", x.value, y.value)
```

Ensuring convexity

- ▶ CVXPY must somehow ensure the written optimization problem is convex. How?
- ▶ **Disciplined convex programming (DCP)**
 - ▶ Defines conventions that ensure an optimization problem is convex
 - ▶ *Example*: the positive sum of two convex functions is convex
 - ▶ These rules are *sufficient* (but not necessary) for convexity
- ▶ Usage in CVXPY
 - ▶ must assess the sign and curvature of `cvxpy.Variable` and `cvx.Parameter` types:
 - ▶ `x.sign`: returns sign of `x`
 - ▶ `x.curvature`: returns the curvature of `x`

Example: `quadratic_minimization.ipynb`

Outline

`scipy.optimize`

CVXPY

Example: `quadratic_minimization.ipynb`

Explore minimization methods minimization

- ▶ Consider minimizing the quadratic function

$$f(x) = \sum_{i=1}^n a_i \cdot (x_i - 1)^2$$

- ▶ *Properties:* convex, smooth, minimum at $x^* = (1, \dots, 1)$
- ▶ Let's compare method performance for:
 1. Well-conditioned (narrow distribution of a_i) v. ill-conditioned (wide distribution of a_i)
 2. Low-dimensional (n small) v. high-dimensional (n large)

scipy.opt function implementation

- ▶ Must define function, and optionally gradient and Hessian

```
def fun(x):  
    return 0.5*sum(np.multiply(quadratic_coeff,\  
                               np.square(np.array(x)-np.ones(np.array(x).size))))  
def fun_grad(x):  
    return np.array(np.multiply(quadratic_coeff,np.array(x)\  
                                -np.ones(np.array(x).size)))  
def fun_hess(x):  
    return np.diag(quadratic_coeff)
```

- ▶ To solve, define initial guess x_0 and invoke a solver with the functions as arguments:

```
res = opt.minimize(fun,x0,method='newton-cg',jac=fun_grad,hess=fun_hess)
```

CVXPY setup

Assume we have already specified:

- ▶ `dimension` (int): number of optimization variable n
- ▶ `quadratic_coeff` (numpy.ndarray): array of a_i

```
import cvxpy as cvx
x = cvx.Variable(dimension)
quadratic_coeff_cvx = cvx.Parameter(dimension, sign='Positive')
quadratic_coeff_cvx.value = quadratic_coeff
obj = cvx.Minimize(0.5 * quadratic_coeff.T * cvx.square(x - 1))
prob = cvx.Problem(obj)
prob.solve()
```

- ▶ Note that the objective has to be explicitly coded in CVXPY objective
- ▶ Cannot use black-box functions!

Method comparison


We will compare:

- ▶ Global, no gradients
 - ▶ `differential_evolution`
 - ▶ *Best performance*: non-convex, low-dimensional. Noise okay!
- ▶ Local, no gradients:
 - ▶ Nelder-Mead
 - ▶ CG with finite-difference Jacobian approximations (CGfd)
 - ▶ *Best performance*: well-conditioned, noise-free, low-dimensional
- ▶ Local, gradients:
 - ▶ CG
 - ▶ *Best performance*: well-conditioned, noise-free. High dimensions okay!
- ▶ Local, gradients and Hessians
 - ▶ `newton-cg`
 - ▶ CVXPY (requires convexity)
 - ▶ *Best performance*: noise-free. Ill-conditioning, high dimensions okay!

Example: [quadratic_minimization.ipynb](#)

Low-dimensional, well-conditioned



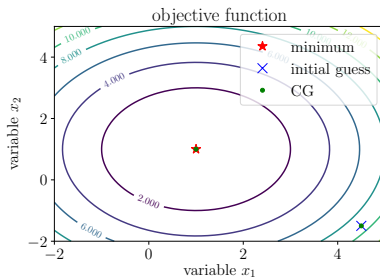
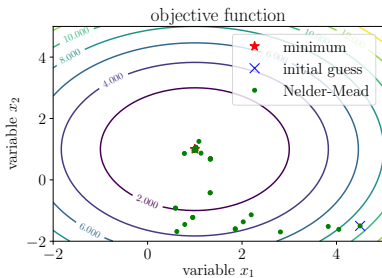
- ▶ Low-dimension: $n = 2$ optimization variables
- ▶ Well-conditioned: $a_i = 1, i = 1, \dots, n$



- ▶ This is the easiest case of all!


Example: `quadratic_minimization.ipynb`

Low-dimensional, well-conditioned



Example: [quadratic_minimization.ipynb](#)

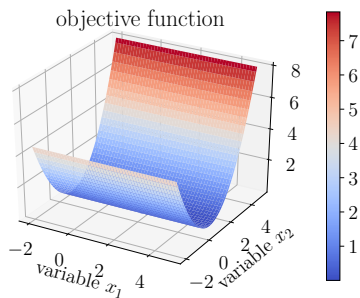
Low-dimensional, well-conditioned



- ▶ All methods find the minimum (computed solution close to $x^* = (1, 1)$)
- ▶ Derivative-free methods (Nelder-Mead and differential evolution) very inefficient!
- ▶ CG more expensive when finite-difference gradient approximations used

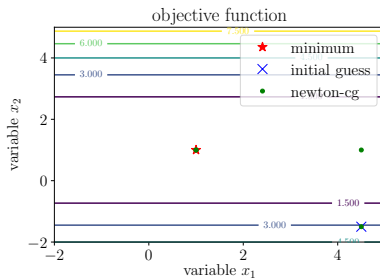



Low-dimensional, poorly conditioned

- ▶ Low-dimension: $n = 2$ optimization variables
- ▶ Poorly conditioned: $a_i = 1$ have large variance ($a_1 = 1.2 \times 10^4$, $a_2 = 1$)




- ▶ Slope is much larger in one direction relative to the other
 - ▶ Hard to minimize in direction x_1 using only the gradient
 - ▶ The Hessian can help in this case!
- Example: [quadratic_minimization.ipynb](#)

Low-dimensional, poorly conditioned



Example: `quadratic_minimization.ipynb`

Low-dimensional, poorly conditioned

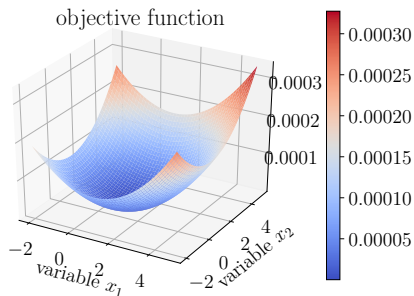


- ▶ All methods do a fairly good job at finding the minimum
- ▶ newton-cg and CVXPY do the best by far (both use Hessian information)
 - ▶ Hessian information helps 'cure' ill conditioning!
- ▶ Derivative-free methods (Nelder-Mead and differential evolution) very inefficient

Example: [quadratic_minimization.ipynb](#)

High-dimensional, poorly conditioned





- ▶ High(er)-dimension: $n = 100$ optimization variables (not truly high dimensional)
- ▶ Poorly conditioned: $a_i = 1$ have large variance ($\max_i a_i / \min_i a_i = 3.6 \times 10^8$)



- ▶ Higher dimensions pose significant challenges to gradient-free methods

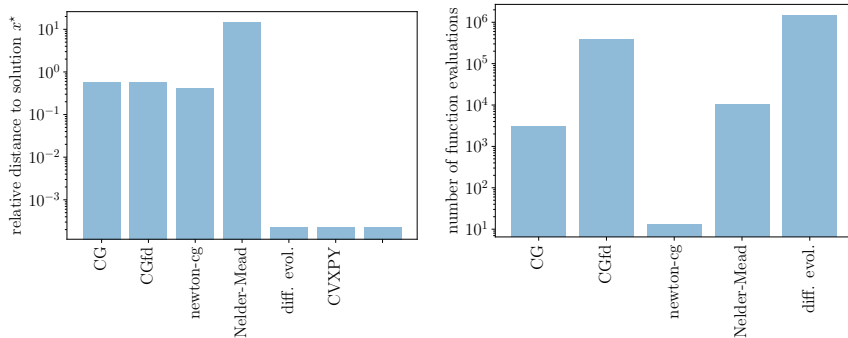
Example: `quadratic_minimization.ipynb`

High-dimensional, poorly conditioned



Example: `quadratic_minimization.ipynb`

High-dimensional, poorly conditioned



- ▶ Nelder–Mead fails to find the minimum in 10,000 function evaluations
- ▶ Differential evolution finds the minimum, but incurs $> 10^6$ function calls!
- ▶ CG w/ finite-difference gradients is very expensive ($n + 1$ function calls per gradient)
- ▶ newton-cg and CVXPY do extremely well (both use Hessian information)

Example: [quadratic_minimization.ipynb](#)

Lessons

- ▶ Gradient information helps “cure” high-dimensionality
 - ▶ Gradients enable a good direction to be found in a high-dimensional space
 - ▶ Without gradients, many function evaluations are needed to explore the space
 - ▶ Finite-difference approximations of the Jacobian become expensive in high dimensions (require $n + 1$ function evaluations)
- ▶ Hessian information helps “cure” ill conditioning!
 - ▶ Hessians inform the optimizer of curvature; thus the optimizer deals with ill conditioning directly
 - ▶ Ill-conditioned Hessians can still pose numerical problems

Let's add noise


- ▶ Let's add sinusoidal noise to the function:

$$f(x) = \sum_{i=1}^n a_i \cdot (x_i - 1)^2 + b \cdot \left[n - \sum_{i=1}^n \cos(2\pi(x_i - 1)) \right]$$

- ▶ b controls the amount of additional noise
- ▶ For $b > 0$, the function is no longer convex!
 - ▶ Many local minima
 - ▶ Local methods may not find the global minimum!
 - ▶ CVXPY not applicable

Low-dimensional, well-conditioned, noisy

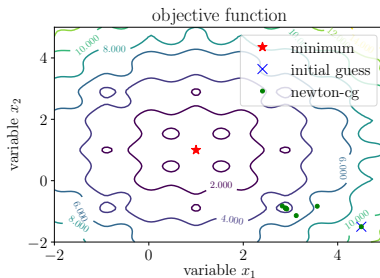

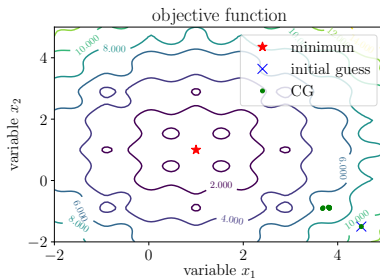
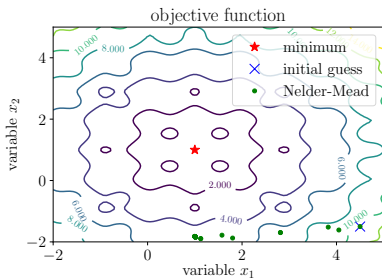
- ▶ Low-dimension: $n = 2$ optimization variables
- ▶ Well-conditioned: $a_i = 1, i = 1, \dots, n$



- ▶ Many local minima in which to get “trapped”


Example: `quadratic_minimization.ipynb`

Low-dimensional, well-conditioned, noisy



Example: `quadratic_minimization.ipynb`


Low-dimensional, well-conditioned, noisy



- ▶ All local methods get trapped in a local minimum
- ▶ CVXPY cannot be used
- ▶ differential evolution finds the closest solution,
 - ▶ However, it requires over a thousand function evaluations!

Example: [quadratic_minimization.ipynb](#)

High-dimensional ($n = 100$), well-conditioned, noisy



- ▶ All local methods get trapped in a local minimum (again)
- ▶ CVXPY cannot be used (again)
- ▶ Differential evolution comes closest to finding the solution
 - ▶ However, it requires over one million function evaluations!

Example: [quadratic_minimization.ipynb](#)

Lessons

Noise can make optimization very difficult!

- ▶ Makes the problem non-convex, with many local minima
- ▶ Local methods get trapped in a local minimum
- ▶ Global methods are needed, but these perform poorly in high dimensions
- ▶ Tools like CVXPY cannot be used
- ▶ **Lesson:** avoid noisy functions by any means possible (e.g., smoothing, convexification)

Recap

- ▶ Global, no gradients
 - ▶ `differential_evolution`
 - ▶ *Best performance*: non-convex, low-dimensional. Noise okay!
- ▶ Local, no gradients:
 - ▶ Nelder-Mead
 - ▶ CG with finite-difference Jacobian approximations (CGfd)
 - ▶ *Best performance*: well-conditioned, noise-free, low-dimensional
- ▶ Local, gradients:
 - ▶ CG
 - ▶ *Best performance*: well-conditioned, noise-free. High dimensions okay!
- ▶ Local, gradients and Hessians
 - ▶ `newton-cg`
 - ▶ CVXPY (requires convexity)
 - ▶ *Best performance*: noise-free. Ill-conditioning, high dimensions okay!